

```
# Data Manipulation and Model Visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import plot_tree

# Supervised Learning

# Regression
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor

# Classification
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB

# Unsupervised Learning
from sklearn.cluster import KMeans # Kmeans

# Data Preprocessing & Feature Engineering
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler

# Dimensionality Reduction
from sklearn.decomposition import PCA

# Model Selection
from sklearn.model_selection import train_test_split

# Evaluation Metrics for Classification
from sklearn import metrics

# Evaluation Metrics for Regression
from sklearn.metrics import root_mean_squared_error,
mean_absolute_error
```

Supervised

Applying Regression

1. Linear Regression

```
# read file
df_linear_regression = pd.read_csv(r"path\Real estate.csv")

# get the first five rows of you data
df_linear_regression.head()

   No  X1 transaction date  X2 house age  \
0    1          2012.917          32.0
1    2          2012.917          19.5
2    3          2013.583          13.3
3    4          2013.500          13.3
4    5          2012.833           5.0

   X3 distance to the nearest MRT station  X4 number of convenience
stores  \
0                                84.87882
10
1                                306.59470
9
2                                561.98450
5
3                                561.98450
5
4                                390.56840
5

   X5 latitude  X6 longitude  Y house price of unit area
0    24.98298    121.54024          37.9
1    24.98034    121.53951          42.2
2    24.98746    121.54391          47.3
3    24.98746    121.54391          54.8
4    24.97937    121.54245          43.1

df_linear_regression.shape # check the shape row & columns of your
data

(414, 8)

df_linear_regression.info() # check basic info of you data

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 414 entries, 0 to 413
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype

```

```

---
0    No                                414 non-null    int64
1    X1 transaction date              414 non-null    float64
2    X2 house age                    414 non-null    float64
3    X3 distance to the nearest MRT station 414 non-null    float64
4    X4 number of convenience stores  414 non-null    int64
5    X5 latitude                     414 non-null    float64
6    X6 longitude                    414 non-null    float64
7    Y house price of unit area       414 non-null    float64
dtypes: float64(6), int64(2)
memory usage: 26.0 KB

```

```
df_linear_regression.isnull().sum() # check the null values of your data
```

```

No                                0
X1 transaction date              0
X2 house age                    0
X3 distance to the nearest MRT station 0
X4 number of convenience stores  0
X5 latitude                     0
X6 longitude                    0
Y house price of unit area       0
dtype: int64

```

```
df_linear_regression.drop(columns=['No'], axis=1, inplace=True) # drop column since i wasn't required
df_linear_regression.head() # again check head of your data
```

```

X1 transaction date  X2 house age  X3 distance to the nearest MRT
station \
0                2012.917         32.0
84.87882
1                2012.917         19.5
306.59470
2                2013.583         13.3
561.98450
3                2013.500         13.3
561.98450
4                2012.833          5.0
390.56840

```

```

X4 number of convenience stores  X5 latitude  X6 longitude \
0                               10    24.98298    121.54024
1                               9     24.98034    121.53951
2                               5     24.98746    121.54391
3                               5     24.98746    121.54391
4                               5     24.97937    121.54245

```

```
Y house price of unit area
```

0	37.9
1	42.2
2	47.3
3	54.8
4	43.1

```
df_linear_regression.describe() # check basic statistics of you data
```

	X1 transaction date	X2 house age \
count	414.000000	414.000000
mean	2013.148971	17.712560
std	0.281967	11.392485
min	2012.667000	0.000000
25%	2012.917000	9.025000
50%	2013.167000	16.100000
75%	2013.417000	28.150000
max	2013.583000	43.800000

	X3 distance to the nearest MRT station \
count	414.000000
mean	1083.885689
std	1262.109595
min	23.382840
25%	289.324800
50%	492.231300
75%	1454.279000
max	6488.021000

	X4 number of convenience stores	X5 latitude	X6 longitude \
count	414.000000	414.000000	414.000000
mean	4.094203	24.969030	121.533361
std	2.945562	0.012410	0.015347
min	0.000000	24.932070	121.473530
25%	1.000000	24.963000	121.528085
50%	4.000000	24.971100	121.538630
75%	6.000000	24.977455	121.543305
max	10.000000	25.014590	121.566270

	Y house price of unit area
count	414.000000
mean	37.980193
std	13.606488
min	7.600000
25%	27.700000
50%	38.450000
75%	46.600000
max	117.500000

```
df_linear_regression.corr() # check the correlation of you data
```

	X1 transaction date	X2 house
age \		
X1 transaction date	1.000000	
0.017549		
X2 house age	0.017549	
1.000000		
X3 distance to the nearest MRT station	0.060880	
0.025622		
X4 number of convenience stores	0.009635	
0.049593		
X5 latitude	0.035058	
0.054420		
X6 longitude	-0.041082	-
0.048520		
Y house price of unit area	0.087491	-
0.210567		

	X3 distance to the nearest MRT
station \	
X1 transaction date	
0.060880	
X2 house age	
0.025622	
X3 distance to the nearest MRT station	
1.000000	
X4 number of convenience stores	-
0.602519	
X5 latitude	-
0.591067	
X6 longitude	-
0.806317	
Y house price of unit area	-
0.673613	

	X4 number of convenience
stores \	
X1 transaction date	
0.009635	
X2 house age	
0.049593	
X3 distance to the nearest MRT station	-
0.602519	
X4 number of convenience stores	
1.000000	
X5 latitude	
0.444143	
X6 longitude	
0.449099	
Y house price of unit area	
0.571005	

	X5 latitude	X6 longitude \
X1 transaction date	0.035058	-0.041082
X2 house age	0.054420	-0.048520
X3 distance to the nearest MRT station	-0.591067	-0.806317
X4 number of convenience stores	0.444143	0.449099
X5 latitude	1.000000	0.412924
X6 longitude	0.412924	1.000000
Y house price of unit area	0.546307	0.523287

	Y house price of unit area
X1 transaction date	0.087491
X2 house age	-0.210567
X3 distance to the nearest MRT station	-0.673613
X4 number of convenience stores	0.571005
X5 latitude	0.546307
X6 longitude	0.523287
Y house price of unit area	1.000000

```
x_linear_regression = df_linear_regression.drop(columns=['Y house price of unit area']) # make independent variables (here drop your target variable)
x_linear_regression.head() # check head of your in independent variables
```

	X1 transaction date	X2 house age	X3 distance to the nearest MRT station \
0	2012.917	32.0	84.87882
1	2012.917	19.5	306.59470
2	2013.583	13.3	561.98450
3	2013.500	13.3	561.98450
4	2012.833	5.0	390.56840

	X4 number of convenience stores	X5 latitude	X6 longitude
0	10	24.98298	121.54024
1	9	24.98034	121.53951
2	5	24.98746	121.54391
3	5	24.98746	121.54391
4	5	24.97937	121.54245

```
y_linear_regression = df_linear_regression['Y house price of unit area'] # make dependent variables (here just select your target variable)
y_linear_regression.head() # check head of your in dependent variable
```

```
0    37.9
1    42.2
2    47.3
3    54.8
4    43.1
```

Name: Y house price of unit area, dtype: float64

```
print("x=",x_linear_regression.shape,"\ny=",
y_linear_regression.shape) # check the shape of you x and y variables
```

```
x= (414, 6)
y= (414,)
```

```
from sklearn.model_selection import train_test_split # import train
test split
```

```
# Apply train_test_split and make your training and testinig x and y
x_train_linear_regression, x_test_linear_regression,
y_train_linear_regression, y_test_linear_regression =
train_test_split(x_linear_regression, y_linear_regression,
test_size=0.3, random_state=42)
```

```
# check the shape of you training and testinig x and y
print("\nx train shape: ", x_train_linear_regression.shape,"\nx test
shape: ", x_test_linear_regression.shape)
print("\ny train shape: ", y_train_linear_regression.shape, "\ny test
shape: ", y_test_linear_regression.shape)
```

```
x train shape: (289, 6)
x test shape: (125, 6)
```

```
y train shape: (289,)
y test shape: (125,)
```

```
from sklearn.linear_model import LinearRegression # import Linear
Regression
```

```
linear_regression = LinearRegression() # call LinearRegression
```

```
linear_regression.fit(x_train_linear_regression,
y_train_linear_regression) # fit you training x and y in
LinearRegression
```

```
LinearRegression()
```

```
linear_regression.coef_ # check the Coedicients
```

```
array([ 5.84779706e+00, -2.42545813e-01, -5.13873381e-03,
1.07449530e+00,
       2.39096949e+02, -5.22351591e+01])
```

```
pd.DataFrame(linear_regression.coef_, x_linear_regression.columns,
columns=["Coedicients"]) # making dataframe of Coedicients
```

	Coedicients
X1 transaction date	5.847797
X2 house age	-0.242546
X3 distance to the nearest MRT station	-0.005139
X4 number of convenience stores	1.074495
X5 latitude	239.096949
X6 longitude	-52.235159

```
y_pred_linear_regression =
linear_regression.predict(x_test_linear_regression) # predict values
by linear regression using x test
y_pred_linear_regression[:5] # just to see first five prrdicted values
(no necessary)
```

```
array([47.55430212, 41.08372744, 44.25551663, 40.51685112,
27.43467608])
```

```
from sklearn import metrics # import metrics
```

```
# check error by placing y test and and y predicted in comparsion
```

```
MAE_linear_regression =
metrics.mean_absolute_error(y_test_linear_regression,
y_pred_linear_regression)
MSE_linear_regression =
metrics.mean_squared_error(y_test_linear_regression,
y_pred_linear_regression)
RMSE_linear_regression =
metrics.root_mean_squared_error(y_test_linear_regression,
y_pred_linear_regression)
```

```
print("Logistic Regression MAE", MAE_linear_regression)
print("Logistic Regression MSE", MSE_linear_regression)
print("Logistic Regression RMSE", RMSE_linear_regression)
```

```
Logistic Regression MAE 6.1848363400971085
Logistic Regression MSE 73.5683793285023
Logistic Regression RMSE 8.577201136064275
```

```
test_residual_linear_regression = y_test_linear_regression -
y_pred_linear_regression # check test residual
```

```
import seaborn as sns # Import seaborn
import matplotlib.pyplot as plt # Import matplotlib
```


```
# Plot 1: Distribution of residuals (displot)
plt.figure(figsize=(6, 6)) # Increase figure size for clarity
sns.displot(test_residual_linear_regression, bins=25, kde=True) #
Displot with KDE
```

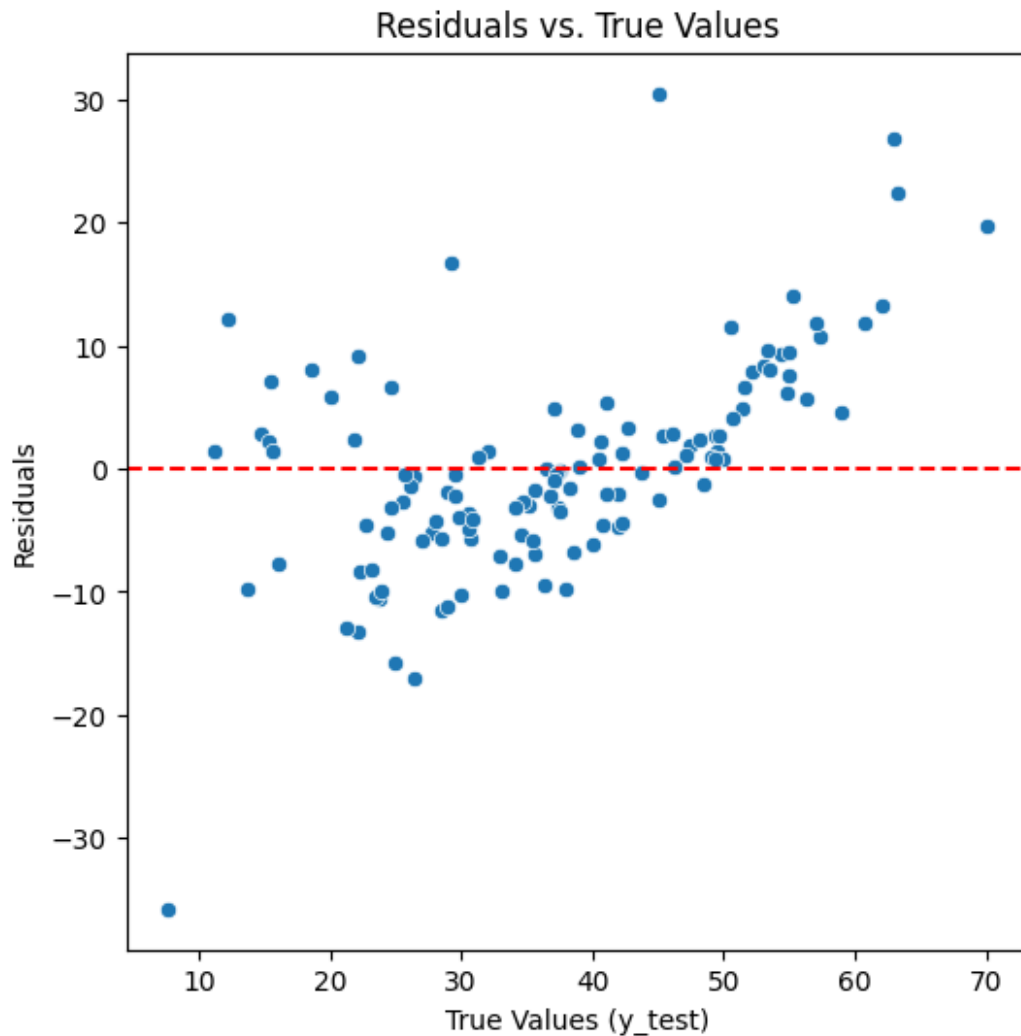


```
plt.title("Distribution of Residuals") # Add title for clarity
plt.xlabel("Residuals") # Label the x-axis
plt.ylabel("Density") # Label the y-axis
plt.show() # Display the plot

# Plot 2: Scatterplot of residuals vs. true values (y_test)
plt.figure(figsize=(6, 6)) # Increase figure size
sns.scatterplot(x=y_test_linear_regression,
y=test_residual_linear_regression) # Scatter plot
plt.axhline(y=0, color='r', ls='--') # Red dashed line at y=0
plt.title("Residuals vs. True Values") # Add title for clarity
plt.xlabel("True Values (y_test)") # Label the x-axis
plt.ylabel("Residuals") # Label the y-axis
plt.show() # Display the plot
```

<Figure size 600x600 with 0 Axes>





2. Decision Tree

```
# read file
df_decision_tree = pd.read_csv(r"path\Real estate.csv")

df_decision_tree.head() # see first five rows
```

	No	X1 transaction date	X2 house age \
0	1	2012.917	32.0
1	2	2012.917	19.5
2	3	2013.583	13.3
3	4	2013.500	13.3
4	5	2012.833	5.0

	X3 distance to the nearest MRT station	X4 number of convenience stores \
0		84.87882

```

10
1          306.59470
9
2          561.98450
5
3          561.98450
5
4          390.56840
5

```

```

      X5 latitude  X6 longitude  Y house price of unit area
0      24.98298      121.54024          37.9
1      24.98034      121.53951          42.2
2      24.98746      121.54391          47.3
3      24.98746      121.54391          54.8
4      24.97937      121.54245          43.1

```

```
df_decision_tree.shape # check the shape (rows and columns)
```

```
(414, 8)
```

```
df_decision_tree.info() # check the basic info
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 414 entries, 0 to 413
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	No	414 non-null	int64
1	X1 transaction date	414 non-null	float64
2	X2 house age	414 non-null	float64
3	X3 distance to the nearest MRT station	414 non-null	float64
4	X4 number of convenience stores	414 non-null	int64
5	X5 latitude	414 non-null	float64
6	X6 longitude	414 non-null	float64
7	Y house price of unit area	414 non-null	float64

```
dtypes: float64(6), int64(2)
```

```
memory usage: 26.0 KB
```

```
df_decision_tree.drop(columns=['No'], axis=1, inplace=True) # drop
`No` bcz its not req.
```

```
df_decision_tree.head() # check head after dropping `No`
```

```

      X1 transaction date  X2 house age  X3 distance to the nearest MRT
station \
0          2012.917          32.0
84.87882
1          2012.917          19.5
306.59470
2          2013.583          13.3
561.98450

```

3	2013.500	13.3
561.98450		
4	2012.833	5.0
390.56840		

	X4 number of convenience stores	X5 latitude	X6 longitude \
0	10	24.98298	121.54024
1	9	24.98034	121.53951
2	5	24.98746	121.54391
3	5	24.98746	121.54391
4	5	24.97937	121.54245

	Y house price of unit area
0	37.9
1	42.2
2	47.3
3	54.8
4	43.1

```
df_decision_tree.isnull().sum() # check any null value
```

X1 transaction date	0
X2 house age	0
X3 distance to the nearest MRT station	0
X4 number of convenience stores	0
X5 latitude	0
X6 longitude	0
Y house price of unit area	0
dtype: int64	

```
df_decision_tree.describe() # check the basic statitics
```

	X1 transaction date	X2 house age \
count	414.000000	414.000000
mean	2013.148971	17.712560
std	0.281967	11.392485
min	2012.667000	0.000000
25%	2012.917000	9.025000
50%	2013.167000	16.100000
75%	2013.417000	28.150000
max	2013.583000	43.800000

	X3 distance to the nearest MRT station \
count	414.000000
mean	1083.885689
std	1262.109595
min	23.382840
25%	289.324800
50%	492.231300
75%	1454.279000

```
max 6488.021000
```

```
      X4 number of convenience stores  X5 latitude  X6 longitude \
count 414.000000 414.000000 414.000000
mean  4.094203 24.969030 121.533361
std    2.945562  0.012410  0.015347
min    0.000000 24.932070 121.473530
25%    1.000000 24.963000 121.528085
50%    4.000000 24.971100 121.538630
75%    6.000000 24.977455 121.543305
max   10.000000 25.014590 121.566270
```

```
      Y house price of unit area
count 414.000000
mean  37.980193
std   13.606488
min    7.600000
25%   27.700000
50%   38.450000
75%   46.600000
max   117.500000
```

```
df_decision_tree.corr() # check the corr relation
```

```
      X1 transaction date  X2 house
age \
X1 transaction date      1.000000
0.017549
X2 house age            0.017549
1.000000
X3 distance to the nearest MRT station  0.060880
0.025622
X4 number of convenience stores        0.009635
0.049593
X5 latitude                    0.035058
0.054420
X6 longitude                   -0.041082  -
0.048520
Y house price of unit area          0.087491  -
0.210567
```

```
      X3 distance to the nearest MRT
station \
X1 transaction date  0.060880
X2 house age        0.025622
X3 distance to the nearest MRT station  1.000000
X4 number of convenience stores  -
```

```

0.602519
X5 latitude -
0.591067
X6 longitude -
0.806317
Y house price of unit area -
0.673613

X4 number of convenience
stores \
X1 transaction date
0.009635
X2 house age
0.049593
X3 distance to the nearest MRT station -
0.602519
X4 number of convenience stores
1.000000
X5 latitude
0.444143
X6 longitude
0.449099
Y house price of unit area
0.571005

X5 latitude X6 longitude \
X1 transaction date 0.035058 -0.041082
X2 house age 0.054420 -0.048520
X3 distance to the nearest MRT station -0.591067 -0.806317
X4 number of convenience stores 0.444143 0.449099
X5 latitude 1.000000 0.412924
X6 longitude 0.412924 1.000000
Y house price of unit area 0.546307 0.523287

Y house price of unit area
X1 transaction date 0.087491
X2 house age -0.210567
X3 distance to the nearest MRT station -0.673613
X4 number of convenience stores 0.571005
X5 latitude 0.546307
X6 longitude 0.523287
Y house price of unit area 1.000000

x_decision_tree = df_decision_tree.drop(columns=['Y house price of
unit area']) # make you independent variables by dropping target variable
in them
x_decision_tree.head() # see the head of your independent variables

X1 transaction date X2 house age X3 distance to the nearest MRT
station \

```

0	2012.917	32.0
84.87882		
1	2012.917	19.5
306.59470		
2	2013.583	13.3
561.98450		
3	2013.500	13.3
561.98450		
4	2012.833	5.0
390.56840		

	X4 number of convenience stores	X5 latitude	X6 longitude
0	10	24.98298	121.54024
1	9	24.98034	121.53951
2	5	24.98746	121.54391
3	5	24.98746	121.54391
4	5	24.97937	121.54245

```
y_decision_tree = df_decision_tree['Y house price of unit area'] #
select your target variable by picking it up this way
y_decision_tree.head() # check the head of ur dep_var.
```

```
0    37.9
1    42.2
2    47.3
3    54.8
4    43.1
```

```
Name: Y house price of unit area, dtype: float64
```

```
from sklearn.model_selection import train_test_split # import
train_test_split
```

```
# split you data in x and y in from of training and testing data
x_train_decision_tree, x_test_decision_tree, y_train_decision_tree,
y_test_decision_tree = train_test_split(x_decision_tree,
y_decision_tree, test_size=0.3, random_state=42)
```

```
# check the shape of your training and testing data
print("Training data shape: ", x_train_decision_tree.shape,
y_train_decision_tree.shape)
print("Testing data shape: ", x_test_decision_tree.shape,
y_test_decision_tree.shape)
```

```
Training data shape: (289, 6) (289,)
Testing data shape: (125, 6) (125,)
```

```
from sklearn.tree import DecisionTreeRegressor # import decision tree
decision_tree = DecisionTreeRegressor(random_state=42) # call decision
tree in a variable
```

```
decision_tree.fit(x_train_decision_tree, y_train_decision_tree) # fit
ur training data
```

```
DecisionTreeRegressor(random_state=42)
```

```
y_pred_decision_tree = decision_tree.predict(x_test_decision_tree) #
predict values by decision tree by using `x_test_decision_tree`
y_pred_decision_tree[:5] # just to see first five predicted values
```

```
array([48.6, 38.9, 43.5, 29.3, 25.7])
```

```
from sklearn import metrics # import metrics
```

```
# check the error value
```

```
MAE_decision_tree = metrics.mean_absolute_error(y_test_decision_tree,
y_pred_decision_tree)
```

```
MSE_decision_tree = metrics.mean_squared_error(y_test_decision_tree,
y_pred_decision_tree)
```

```
RMSE_decision_tree =
metrics.root_mean_squared_error(y_test_decision_tree,
y_pred_decision_tree)
```

```
# print error values
```

```
print("Decision Tree MAE", MAE_decision_tree)
```

```
print("Decision Tree MSE", MSE_decision_tree)
```

```
print("Decision Tree RMSE", RMSE_decision_tree)
```

```
Decision Tree MAE 5.6076
```

```
Decision Tree MSE 73.4509
```

```
Decision Tree RMSE 8.570350051193943
```

```
test_residual_decision_tree = y_test_decision_tree -
y_pred_decision_tree
test_residual_decision_tree
```

```
358    -3.5
```

```
350     3.4
```

```
373     8.7
```

```
399     8.0
```

```
369    -2.9
```

```
...
```

```
268    -0.5
```

```
148    26.9
```

```
16     0.4
```

```
66    -3.7
```

```
341    -9.4
```

```
Name: Y house price of unit area, Length: 125, dtype: float64
```

```
import seaborn as sns # Import seaborn
```

```
import matplotlib.pyplot as plt # Import matplotlib
```




```


# Plot 1: Distribution of residuals (displot)
plt.figure(figsize=(6, 6)) # Increase figure size for clarity
sns.displot(test_residual_decision_tree, bins=25, kde=True) # Displot
with KDE
plt.title("Distribution of Residuals") # Add title for clarity
plt.xlabel("Residuals") # Label the x-axis
plt.ylabel("Density") # Label the y-axis
plt.show() # Display the plot

# Plot 2: Scatterplot of residuals vs. true values (y_test)
plt.figure(figsize=(6, 6)) # Increase figure size
sns.scatterplot(x=y_test_decision_tree, y=test_residual_decision_tree)
# Scatter plot
plt.axhline(y=0, color='r', ls='--') # Red dashed line at y=0
plt.title("Residuals vs. True Values") # Add title for clarity
plt.xlabel("True Values (y_test)") # Label the x-axis
plt.ylabel("Residuals") # Label the y-axis
plt.show() # Display the plot

```

<Figure size 600x600 with 0 Axes>





things are approxly same I will not use comment and and not repeat same steps

3. Random Foest

```
df_random_forest = pd.read_csv(r"path\Real estate.csv")
df_random_forest.head()
```

	No	X1 transaction date	X2 house age \
0	1	2012.917	32.0
1	2	2012.917	19.5
2	3	2013.583	13.3
3	4	2013.500	13.3
4	5	2012.833	5.0

	X3 distance to the nearest MRT station	X4 number of convenience stores \
0		84.87882
10		

1	306.59470
9	
2	561.98450
5	
3	561.98450
5	
4	390.56840
5	

	X5 latitude	X6 longitude	Y house price of unit area
0	24.98298	121.54024	37.9
1	24.98034	121.53951	42.2
2	24.98746	121.54391	47.3
3	24.98746	121.54391	54.8
4	24.97937	121.54245	43.1

```
df_random_forest.drop(columns=['No'], axis=True, inplace=True)
x_random_forest = df_random_forest.drop(columns=['Y house price of
unit area'])
y_random_forest = df_random_forest['Y house price of unit area']

x_train_random_forest, x_test_random_forest, y_train_random_forest,
y_test_random_forest = train_test_split(x_random_forest,
y_random_forest, test_size=0.3, random_state=42)
```

```
random_forest = RandomForestRegressor(random_state=42)
random_forest.fit(x_train_random_forest, y_train_random_forest)
```

```
RandomForestRegressor(random_state=42)
```

```
y_pred_random_forest = random_forest.predict(x_test_random_forest)
y_pred_random_forest[:5]
```

```
array([49.153 , 38.522 , 53.18325, 35.144 , 25.453 ])
```

```
print("MAE Random Froest",
metrics.mean_absolute_error(y_test_random_forest,
y_pred_random_forest))
print("MSE Random Froest",
metrics.mean_squared_error(y_test_random_forest,
y_pred_random_forest))
print("RMSE Random Froest",
metrics.root_mean_squared_error(y_test_random_forest,
y_pred_random_forest))
```

```
MAE Random Froest 4.397113295238097
MSE Random Froest 42.75048519970811
RMSE Random Froest 6.538385519354767
```

```

test_residual_random_forest = y_test_random_forest -
y_pred_random_forest
test_residual_random_forest

358      -4.05300
350       3.77800
373      -0.98325
399       2.15600
369      -2.65300
...
268      -2.35360
148      19.87900
16       2.31025
66       -1.42950
341      -1.53900
Name: Y house price of unit area, Length: 125, dtype: float64

import seaborn as sns # Import seaborn
import matplotlib.pyplot as plt # Import matplotlib


# Plot 1: Distribution of residuals (displot)
plt.figure(figsize=(6, 6)) # Increase figure size for clarity
sns.displot(test_residual_random_forest, bins=25, kde=True) # Displot
with KDE
plt.title("Distribution of Residuals") # Add title for clarity
plt.xlabel("Residuals") # Label the x-axis
plt.ylabel("Density") # Label the y-axis
plt.show() # Display the plot


# Plot 2: Scatterplot of residuals vs. true values (y_test)
plt.figure(figsize=(6, 6)) # Increase figure size
sns.scatterplot(x=y_pred_random_forest, y=test_residual_random_forest)
# Scatter plot
plt.axhline(y=0, color='r', ls='--') # Red dashed line at y=0
plt.title("Residuals vs. True Values") # Add title for clarity
plt.xlabel("True Values (y_test)") # Label the x-axis
plt.ylabel("Residuals") # Label the y-axis
plt.show() # Display the plot

<Figure size 600x600 with 0 Axes>

```

Distribution of Residuals





4. SVM

```
df_svm = pd.read_csv(r"path\Real estate.csv")
df_svm.head()
```

	No	X1 transaction date	X2 house age \
0	1	2012.917	32.0
1	2	2012.917	19.5
2	3	2013.583	13.3
3	4	2013.500	13.3
4	5	2012.833	5.0

	X3 distance to the nearest MRT station	X4 number of convenience stores \
0		84.87882
10		
1		306.59470

0	84.87882
10	
1	306.59470

```

9
2
5
3
5
4
5

```

	X5 latitude	X6 longitude	Y house price of unit area
0	24.98298	121.54024	37.9
1	24.98034	121.53951	42.2
2	24.98746	121.54391	47.3
3	24.98746	121.54391	54.8
4	24.97937	121.54245	43.1

```

df_svm.drop(columns=['No'], axis=1, inplace=True)

x_svm = df_svm.drop(columns=['Y house price of unit area'])
y_svm = df_svm['Y house price of unit area']

x_train_svm, x_test_svm, y_train_svm, y_test_svm =
train_test_split(x_svm, y_svm, test_size=0.3, random_state=42)

svm = SVR()
svm.fit(x_train_svm, y_train_svm)

SVR()

y_pred_svm = svm.predict(x_test_svm)
y_pred_svm[:5]

array([43.90987572, 41.32920139, 43.27836755, 44.0793227 ,
25.17732503])

print("MAE SVM: ", metrics.mean_absolute_error(y_test_svm,
y_pred_svm))
print("MSE SVM: ", metrics.mean_squared_error(y_test_svm, y_pred_svm))
print("RMSE SVM: ", metrics.root_mean_squared_error(y_test_svm,
y_pred_svm))

MAE SVM:  6.882396075217186
MSE SVM:  85.34235717951594
RMSE SVM:  9.23809272412417

test_residual_svm = y_test_svm - y_pred_svm
test_residual_svm

358    1.190124
350    0.970799
373    8.921632
399   -6.779323
369   -2.377325

```

```
...
268    -2.171528
148    24.718366
16     26.977536
66     6.790088
341    -1.707198
```


Name: Y house price of unit area, Length: 125, dtype: float64


```
plt.figure(figsize=(6, 6))
sns.displot(test_residual_random_forest, bins=25, kde=True)
plt.title("Distribution of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Density")
plt.show()

plt.figure(figsize=(6, 6))
sns.scatterplot(x=y_pred_random_forest, y=test_residual_random_forest)
plt.axhline(y=0, color='r', ls='--')
plt.title("Residuals vs. True Values")
plt.xlabel("True Values (y_test)")
plt.ylabel("Residuals")
plt.show()
```

<Figure size 600x600 with 0 Axes>

Distribution of Residuals





5. KNN

```
df_knn = pd.read_csv(r"path\Real estate.csv")
df_knn.head()
```

	No	X1 transaction date	X2 house age \
0	1	2012.917	32.0
1	2	2012.917	19.5
2	3	2013.583	13.3
3	4	2013.500	13.3
4	5	2012.833	5.0

	X3 distance to the nearest MRT station	X4 number of convenience stores \
0		84.87882
10		
1		306.59470

```

9
2
5
3
5
4
5

```

	X5 latitude	X6 longitude	Y house price of unit area
0	24.98298	121.54024	37.9
1	24.98034	121.53951	42.2
2	24.98746	121.54391	47.3
3	24.98746	121.54391	54.8
4	24.97937	121.54245	43.1

```

df_knn.drop(columns=['No'], axis=1, inplace=True)
x_knn = df_knn.drop(columns=['Y house price of unit area'])
y_knn = df_knn['Y house price of unit area']

x_train_knn, x_test_knn, y_train_knn, y_test_knn =
train_test_split(x_knn, y_knn, test_size=0.3, random_state=42)

```

```

knn = KNeighborsRegressor()
knn.fit(x_train_knn, y_train_knn)

```

```

KNeighborsRegressor()

```

```

y_pred_knn = knn.predict(x_test_knn)
y_pred_knn[:5]

```

```

array([49.82, 39.16, 46.06, 38.98, 26.3 ])

```

```

print("MAE KNN: ", metrics.mean_absolute_error(y_test_knn,
y_pred_knn))
print("MSE KNN: ", metrics.mean_squared_error(y_test_knn, y_pred_knn))
print("RMSE KNN: ", metrics.root_mean_squared_error(y_test_knn,
y_pred_knn))

```

```

MAE KNN:  5.45248
MSE KNN:  62.8123712
RMSE KNN:  7.925425616331277

```

```

test_residual_knn = y_test_knn - y_pred_knn
test_residual_knn

```

```

358    -4.72
350     3.14
373     6.14
399    -1.68
369    -3.50
...

```

```
268      1.72
148      22.74
16       5.68
66      -0.46
341     -5.56
```


```
Name: Y house price of unit area, Length: 125, dtype: float64
```

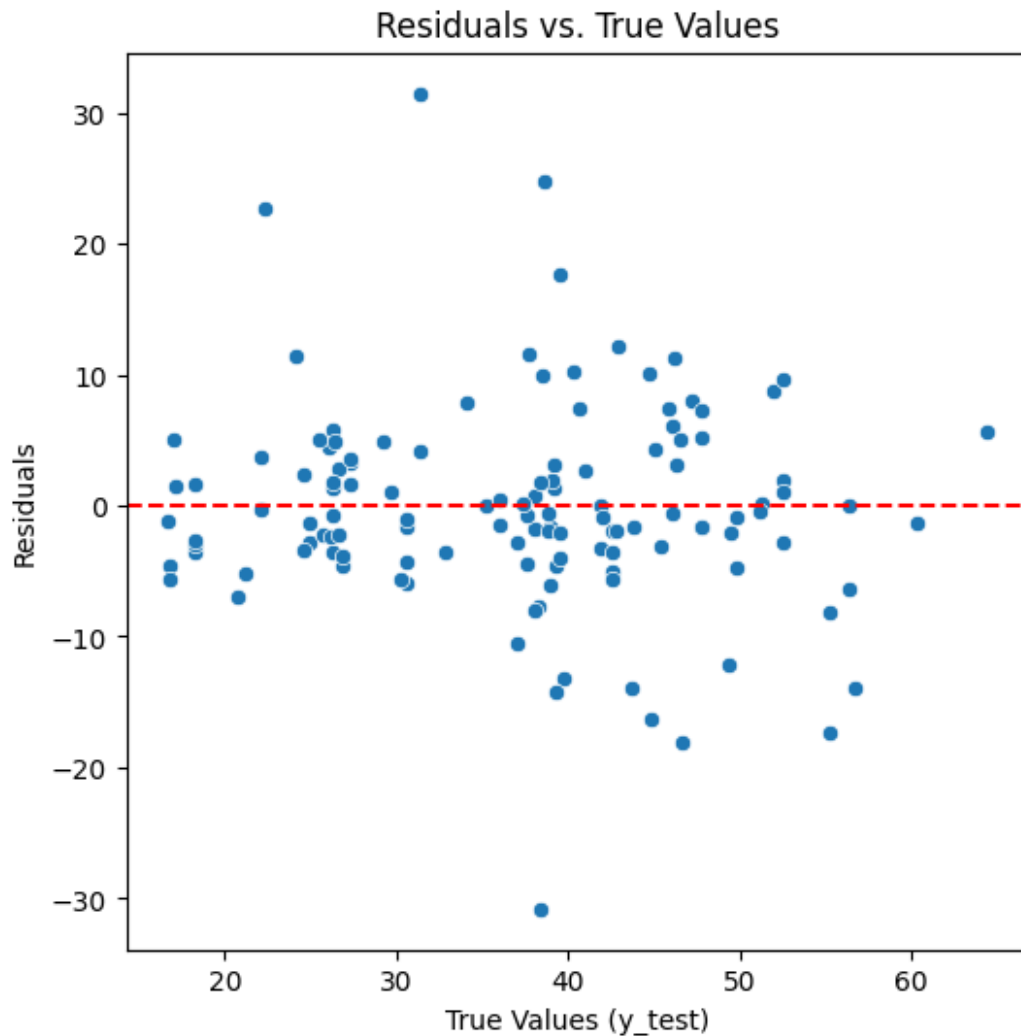
```
plt.figure(figsize=(6, 6))
sns.displot(test_residual_knn, bins=25, kde=True)
plt.title("Distribution of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Density")
plt.show()
```

```
plt.figure(figsize=(6, 6))
sns.scatterplot(x=y_pred_knn, y=test_residual_knn)
plt.axhline(y=0, color='r', ls='--')
plt.title("Residuals vs. True Values")
plt.xlabel("True Values (y_test)")
plt.ylabel("Residuals")
plt.show()
```

```
<Figure size 600x600 with 0 Axes>
```

Distribution of Residuals





Comparison

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

# Create a dictionary with the results for each model and each metric
data = {
    'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest',
    'SVM', 'KNN'],
    'MAE': [6.1848363400971085, 5.6076, 4.397113295238097,
    6.882396075217186, 5.45248],
    'MSE': [73.5683793285023, 73.4509, 42.75048519970811,
    85.34235717951594, 62.8123712],
    'RMSE': [8.577201136064275, 8.570350051193943, 6.538385519354767,
```

```

9.23809272412417, 7.925425616331277]
}

# Convert the dictionary to a pandas DataFrame
df = pd.DataFrame(data)

# Create a figure and axis with 3 subplots (1 row, 3 columns)
fig, axes = plt.subplots(1, 3, figsize=(14, 4))

# Plot MAE, MSE, and RMSE in separate subplots

# MAE subplot
axes[0].bar(df['Model'], df['MAE'], color='b', edgecolor='grey')
axes[0].set_title('Mean Absolute Error (MAE)', fontsize=14)
axes[0].set_ylabel('MAE Value', fontsize=12)
axes[0].tick_params(axis='x', rotation=45) # Rotate x-axis labels for
better readability


# MSE subplot
axes[1].bar(df['Model'], df['MSE'], color='g', edgecolor='grey')
axes[1].set_title('Mean Squared Error (MSE)', fontsize=14)
axes[1].set_ylabel('MSE Value', fontsize=12)
axes[1].tick_params(axis='x', rotation=45)

# RMSE subplot
axes[2].bar(df['Model'], df['RMSE'], color='r', edgecolor='grey')
axes[2].set_title('Root Mean Squared Error (RMSE)', fontsize=14)
axes[2].set_ylabel('RMSE Value', fontsize=12)
axes[2].tick_params(axis='x', rotation=45)

# Adjust layout for better spacing
plt.tight_layout()

# Show the plot
plt.show()

```



Applying Classification

Making one df and x and y and training and testing data

1. Logistic Regression

```
import pandas as pd
df = pd.read_csv(r"path\bank.csv")
df.head()
```

	age	job	marital	education	default	balance	housing	loan
contact \								
0	59	admin.	married	secondary	no	2343	yes	no
unknown								
1	56	admin.	married	secondary	no	45	no	no
unknown								
2	41	technician	married	secondary	no	1270	yes	no
unknown								
3	55	services	married	secondary	no	2476	yes	no
unknown								
4	54	admin.	married	tertiary	no	184	no	no
unknown								

	day	month	duration	campaign	pdays	previous	poutcome	deposit
0	5	may	1042	1	-1	0	unknown	yes
1	5	may	1467	1	-1	0	unknown	yes
2	5	may	1389	1	-1	0	unknown	yes
3	5	may	579	1	-1	0	unknown	yes
4	5	may	673	2	-1	0	unknown	yes

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11162 entries, 0 to 11161
Data columns (total 17 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             11162 non-null  int64
1   job             11162 non-null  object
2   marital         11162 non-null  object
3   education       11162 non-null  object
4   default         11162 non-null  object
5   balance         11162 non-null  int64
6   housing         11162 non-null  object
7   loan           11162 non-null  object
8   contact         11162 non-null  object
9   day             11162 non-null  int64
10  month           11162 non-null  object
11  duration        11162 non-null  int64
```



```

12  campaign    11162 non-null  int64
13  pdays      11162 non-null  int64
14  previous    11162 non-null  int64
15  poutcome    11162 non-null  object
16  deposit     11162 non-null  object

```

```
dtypes: int64(7), object(10)
```

```
memory usage: 1.4+ MB
```

```
df.isnull().sum()
```

```

age          0
job          0
marital      0
education    0
default      0
balance      0
housing      0
loan         0
contact      0
day          0
month        0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
deposit      0
dtype: int64

```

```
categorical_columns = df.select_dtypes(object).columns
```

```
label_encoder = LabelEncoder()
```

```
df1 = df.copy()
```

```

for column in categorical_columns:
    df1[column] = label_encoder.fit_transform(df[column])

```

```
df1
```

	age	job	marital	education	default	balance	housing	loan
contact \								
0	59	0	1	1	0	2343	1	0
2								
1	56	0	1	1	0	45	0	0
2								
2	41	9	1	1	0	1270	1	0
2								
3	55	7	1	1	0	2476	1	0
2								
4	54	0	1	2	0	184	0	0

```

2
...    ...    ...    ...    ...    ...    ...    ...    ...
...
11157  33    1      2          0      0      1      1      0
0
11158  39    7      1          1      0      733     0      0
2
11159  32    9      2          1      0      29      0      0
0
11160  43    9      1          1      0      0       0      1
0
11161  34    9      1          1      0      0       0      0
0

```

```

      day month duration campaign pdays previous poutcome
deposit
0      5     8    1042         1     -1         0         3
1
1      5     8    1467         1     -1         0         3
1
2      5     8    1389         1     -1         0         3
1
3      5     8     579         1     -1         0         3
1
4      5     8     673         2     -1         0         3
1
...    ...    ...    ...    ...    ...    ...    ...
...
11157  20     0    257         1     -1         0         3
0
11158  16     6     83         4     -1         0         3
0
11159  19     1    156         2     -1         0         3
0
11160   8     8     9         2    172         5         0
0
11161   9     5    628         1     -1         0         3
0

```

```
[11162 rows x 17 columns]
```

```

x = df1.drop(columns=['deposit'])
y = df1['deposit']

```

```

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.3, random_state=42)

```

```
logistic_regression = LogisticRegression()
```

```
logistic_regression.fit(x_train, y_train)
```

```
C:\Users\Dell\AppData\Roaming\Python\Python313\site-packages\sklearn\
linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to
converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
LogisticRegression()

y_pred_logistic_regression = logistic_regression.predict(x_test)
y_pred_logistic_regression[:5]

array([0, 1, 1, 1, 0])

print("Logistic Regression Classification Report: \n",
metrics.classification_report(y_test, y_pred_logistic_regression))
print("\n\nLogistic Regression Confusion Matrix: \n",
metrics.confusion_matrix(y_test, y_pred_logistic_regression))
print("\n\n\nLogistic Regression Accuracy Score: \n",
metrics.accuracy_score(y_test, y_pred_logistic_regression))
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.74	0.80	0.77	1742
1	0.76	0.69	0.72	1607
accuracy			0.75	3349
macro avg	0.75	0.75	0.75	3349
weighted avg	0.75	0.75	0.75	3349

Logistic Regression Confusion Matrix:

```
[[1393  349]
 [ 496 1111]]
```

Logistic Regression Accuracy Score:

```
0.7476858763810093
```

2. Decision Tree

```
decision_tree = DecisionTreeClassifier()

decision_tree.fit(x_train, y_train)

DecisionTreeClassifier()

y_pred_decision_tree = decision_tree.predict(x_test)
y_pred_decision_tree[:5]

array([0, 1, 0, 1, 0])

print("Decision Tree Classification Report: \n",
      metrics.classification_report(y_test, y_pred_decision_tree))
print("\n\nDecision Tree Confusion Matrix: \n",
      metrics.confusion_matrix(y_test, y_pred_decision_tree))
print("\n\nDecision Tree Accuracy Score: \n",
      metrics.accuracy_score(y_test, y_pred_decision_tree))
```

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.78	0.79	0.79	1742
1	0.77	0.76	0.77	1607
accuracy			0.78	3349
macro avg	0.78	0.78	0.78	3349
weighted avg	0.78	0.78	0.78	3349

Decision Tree Confusion Matrix:

```
[[1379  363]
 [ 380 1227]]
```

Decision Tree Accuracy Score:

```
0.7781427291728874
```

3. Random Forest

```
random_forest = RandomForestClassifier()

random_forest.fit(x_train, y_train)

RandomForestClassifier()

y_pred_random_forest = random_forest.predict(x_test)
y_pred_random_forest[:5]
```

```
array([1, 1, 1, 1, 0])
```

```
print("Random Forest Classification Report: \n",
metrics.classification_report(y_test, y_pred_random_forest))
print("\n\nRandom Forest Confusion Matrix: \n",
metrics.confusion_matrix(y_test, y_pred_random_forest))
print("\n\nRandom Forest Accuracy Score: \n",
metrics.accuracy_score(y_test, y_pred_random_forest))
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.87	0.81	0.84	1742
1	0.81	0.87	0.84	1607
accuracy			0.84	3349
macro avg	0.84	0.84	0.84	3349
weighted avg	0.84	0.84	0.84	3349

Random Forest Confusion Matrix:

```
[[1418  324]
 [ 213 1394]]
```

Random Forest Accuracy Score:

```
0.8396536279486414
```

4. SVM

```
svm = RandomForestClassifier()
```

```
svm.fit(x_train, y_train)
```

```
RandomForestClassifier()
```

```
y_pred_svm = svm.predict(x_test)
```

```
y_pred_svm[:5]
```

```
array([1, 1, 1, 1, 0])
```

```
print("SVM Classification Report: \n",
metrics.classification_report(y_test, y_pred_svm))
print("\n\nSVM Confusion Matrix: \n", metrics.confusion_matrix(y_test,
y_pred_svm))
print("\n\nSVM Accuracy Score: \n", metrics.accuracy_score(y_test,
y_pred_svm))
```

```
SVM Classification Report:
              precision    recall  f1-score   support

     0       0.87       0.82       0.84       1742
     1       0.81       0.87       0.84       1607

 accuracy          0.84
 macro avg         0.84
weighted avg         0.84

SVM Confusion Matrix:
[[1423  319]
 [ 214 1393]]

SVM Accuracy Score:
0.8408480143326366
```

6. KNN

```
knn = KNeighborsClassifier()
knn.fit(x_train, y_train)
KNeighborsClassifier()
y_pred_knn = knn.predict(x_test)
y_pred_knn[:5]
array([0, 1, 1, 1, 0])

print("KNN Classification Report: \n",
      metrics.classification_report(y_test, y_pred_knn))
print("\n\nKNN Confusion Matrix: \n", metrics.confusion_matrix(y_test,
y_pred_knn))
print("\n\nKNN Accuracy Score: \n", metrics.accuracy_score(y_test,
y_pred_knn))

KNN Classification Report:
              precision    recall  f1-score   support

     0       0.75       0.76       0.76       1742
     1       0.74       0.72       0.73       1607

 accuracy          0.74
 macro avg         0.74
weighted avg         0.74
```

```
KNN Confusion Matrix:
[[1330  412]
 [ 445 1162]]
```

```
KNN Accuracy Score:
0.7441027172290235
```

7. Navie Bayes

```
nb = GaussianNB()
nb.fit(x_train, y_train)
GaussianNB()
y_pred_nb = nb.predict(x_test)
y_pred_nb[:5]
array([1, 1, 1, 1, 0])
print("NB Classification Report: \n",
      metrics.classification_report(y_test, y_pred_nb))
print("\n\nNB Confusion Matrix: \n", metrics.confusion_matrix(y_test,
y_pred_nb))
print("\n\nNB Accuracy Score: \n", metrics.accuracy_score(y_test,
y_pred_nb))
```

```
NB Classification Report:
              precision    recall  f1-score   support

     0       0.78       0.72       0.75       1742
     1       0.72       0.78       0.75       1607

 accuracy          0.75
 macro avg         0.75
weighted avg         0.75
```

```
NB Confusion Matrix:
[[1256  486]
 [ 353 1254]]
```

```
NB Accuracy Score:
0.7494774559570021
```

Comparison

```
# Sample classification report strings for each model
classification_reports = {
    'Logistic Regression': {
        'precision': [0.74, 0.76],
        'recall': [0.80, 0.69],
        'f1-score': [0.77, 0.72],
    },
    'Decision Tree': {
        'precision': [0.79, 0.78],
        'recall': [0.80, 0.77],
        'f1-score': [0.79, 0.77],
    },
    'Random Forest': {
        'precision': [0.87, 0.81],
        'recall': [0.82, 0.87],
        'f1-score': [0.84, 0.84],
    },
    'SVM': {
        'precision': [0.87, 0.81],
        'recall': [0.81, 0.87],
        'f1-score': [0.84, 0.84],
    },
    'KNN': {
        'precision': [0.75, 0.74],
        'recall': [0.76, 0.72],
        'f1-score': [0.76, 0.73],
    },
    'Naive Bayes': {
        'precision': [0.78, 0.72],
        'recall': [0.72, 0.78],
        'f1-score': [0.75, 0.75],
    },
}

# Create a figure with 3 rows and 2 columns (for precision, recall, f1-score)
fig, axes = plt.subplots(3, 2, figsize=(15, 18))

# Define the metrics labels (precision, recall, f1-score)
metrics = ['precision', 'recall', 'f1-score']

# Loop through each metric and plot the heatmap for each model
for i, metric in enumerate(metrics):
    # Prepare a list of model names and corresponding metric values
    values = []
```



```

for model in classification_reports:
    values.append(classification_reports[model][metric])

    # Convert the values into a DataFrame
    df = pd.DataFrame(values, columns=['Class 0', 'Class 1'],
index=classification_reports.keys())

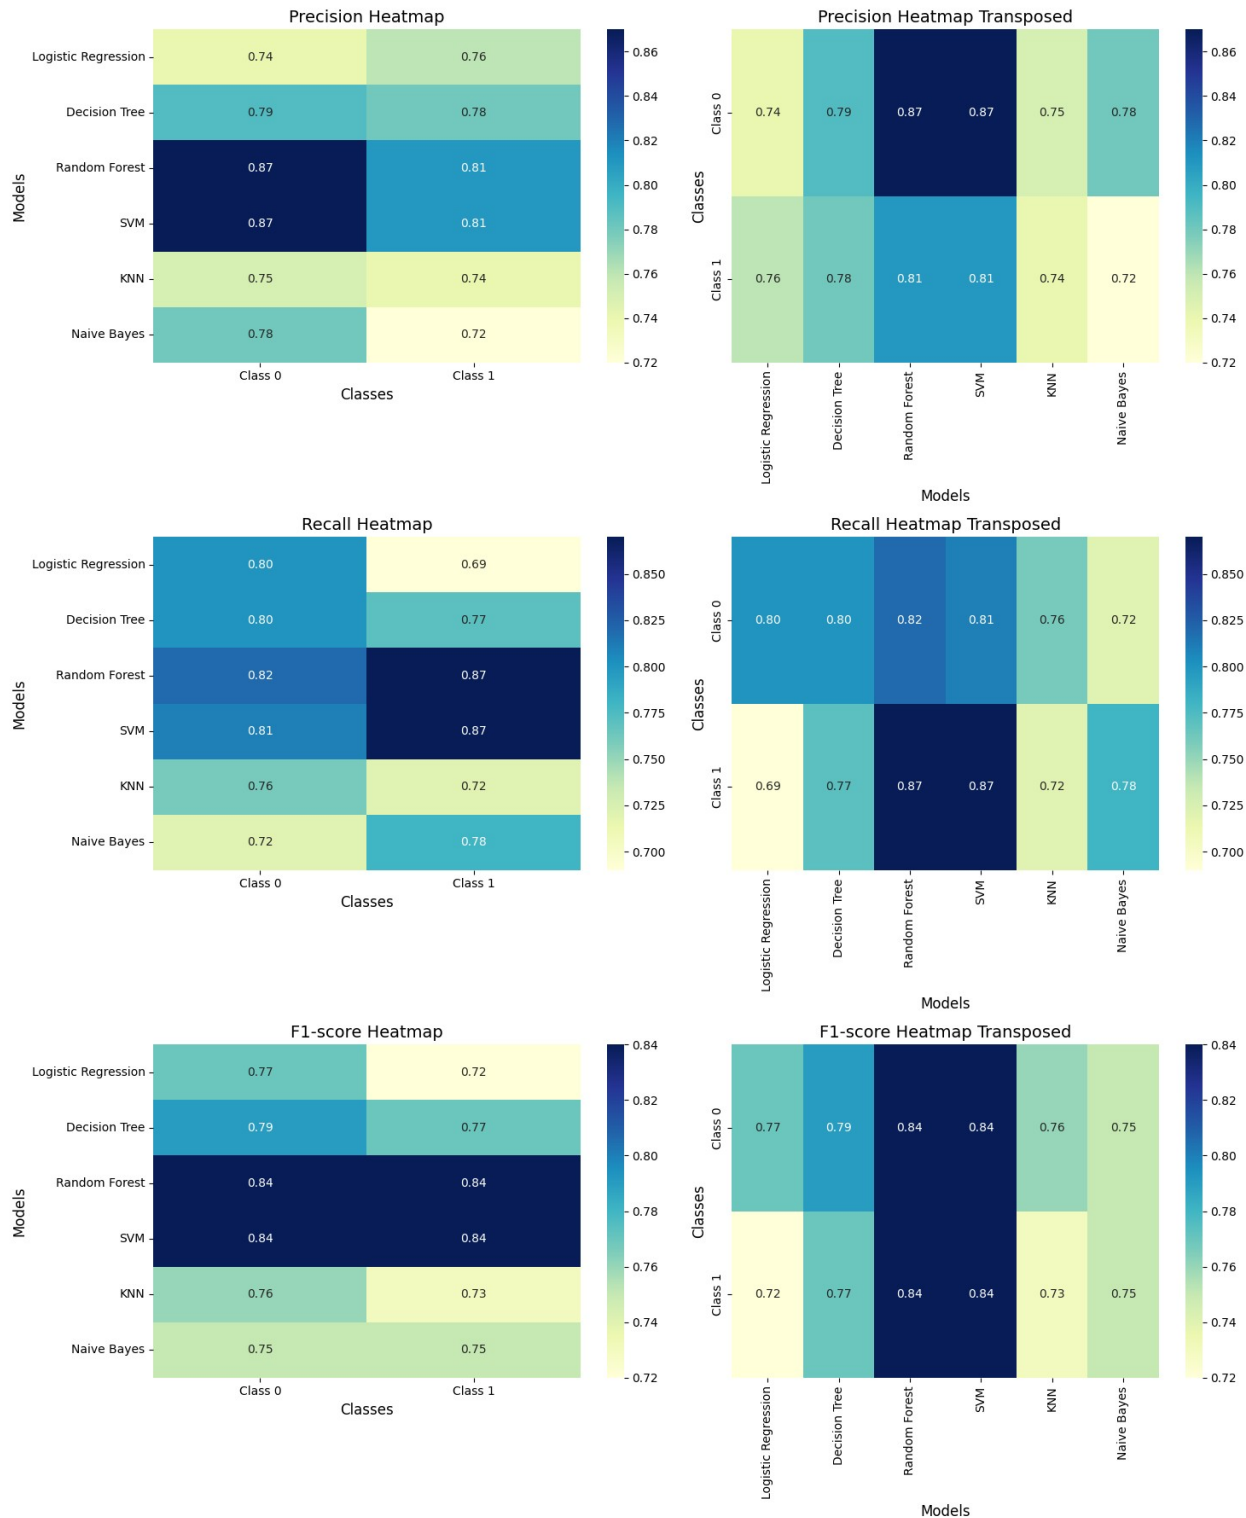
    # Plot the heatmap in the corresponding subplot
    sns.heatmap(df, annot=True, cmap="YlGnBu", cbar=True, fmt=".2f",
ax=axes[i][0])
    axes[i][0].set_title(f'{metric.capitalize()} Heatmap',
fontsize=14)
    axes[i][0].set_ylabel('Models', fontsize=12)
    axes[i][0].set_xlabel('Classes', fontsize=12)

    # Now move to next column (to keep the layout of each model in a
new column)
    sns.heatmap(df.T, annot=True, cmap="YlGnBu", cbar=True, fmt=".2f",
ax=axes[i][1])
    axes[i][1].set_title(f'{metric.capitalize()} Heatmap Transposed',
fontsize=14)
    axes[i][1].set_ylabel('Classes', fontsize=12)
    axes[i][1].set_xlabel('Models', fontsize=12)

# Adjust layout for better spacing
plt.tight_layout()

# Show the plot
plt.show()

```



```
# Confusion matrices for each model
confusion_matrices = {
    'Logistic Regression': np.array([[1393, 349], [496, 1111]]),
    'Decision Tree': np.array([[1390, 352], [368, 1239]]),
```

```

    'Random Forest': np.array([[1421, 321], [216, 1391]]),
    'SVM': np.array([[1419, 323], [216, 1391]]),
    'KNN': np.array([[1330, 412], [445, 1162]]),
    'Naive Bayes': np.array([[1256, 486], [353, 1254]])
}

# Create a figure with 3 rows and 2 columns (for 6 models)
fig, axes = plt.subplots(3, 2, figsize=(15, 18))

# Loop through each confusion matrix and plot the heatmap
for i, (model, matrix) in enumerate(confusion_matrices.items()):
    row, col = divmod(i, 2) # Determine the row and column for subplot

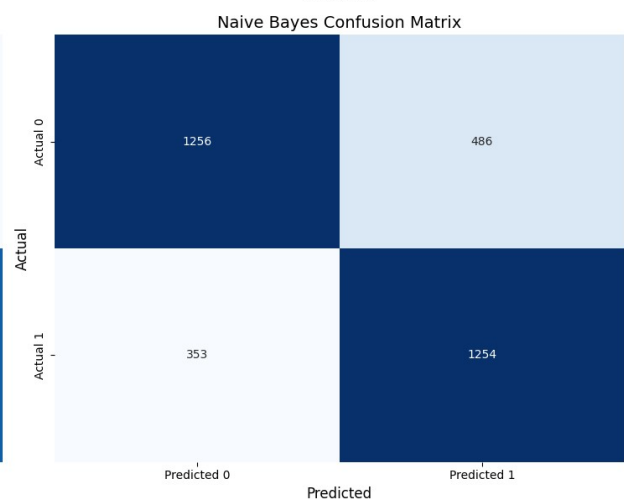




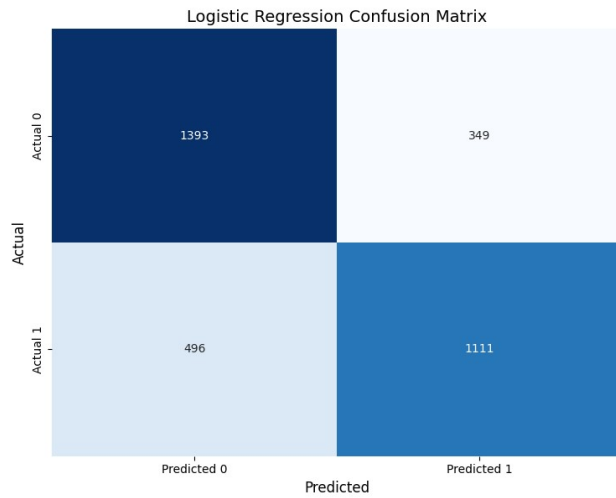
    # Plot the heatmap for the confusion matrix
    sns.heatmap(matrix, annot=True, fmt='d', cmap="Blues", cbar=False,
ax=axes[row][col],
                xticklabels=['Predicted 0', 'Predicted 1'],
yticklabels=['Actual 0', 'Actual 1'])

    # Set the title and labels for the subplot
    axes[row][col].set_title(f'{model} Confusion Matrix', fontsize=14)
    axes[row][col].set_xlabel('Predicted', fontsize=12)
    axes[row][col].set_ylabel('Actual', fontsize=12)

# Adjust layout for better spacing
plt.tight_layout()

# Show the plot
plt.show()

```



```
# Data: Models and their corresponding accuracy scores
models = ['Logistic Regression', 'Decision Tree', 'Random Forest',
'SVM', 'KNN', 'Naive Bayes']
accuracy_scores = [0.7476858763810093, 0.78501045088086,
0.8396536279486414,
```

```
0.8390564347566438, 0.7441027172290235,  
0.7494774559570021]
```

```
# Create a bar chart
```

```
plt.figure(figsize=(14, 5))
```

```
sns.barplot(x=models, y=accuracy_scores, palette='viridis')
```

```
# Set titles and labels
```

```
plt.title('Model Accuracy Comparison', fontsize=16)
```

```
plt.xlabel('Models', fontsize=14)
```

```
plt.ylabel('Accuracy Score', fontsize=14)
```

```
# Rotate x-axis labels for better visibility
```

```
plt.xticks(rotation=45, ha='right')
```

```
# Display the plot
```

```
plt.tight_layout()
```


```
plt.show()
```

C:\Users\Dell\AppData\Local\Temp\ipykernel_19280\2162836733.py:8:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=models, y=accuracy_scores, palette='viridis')
```



THE END @:-)
